



1 B. SQL

Slides adapted from John
Dowell

Structured Query Language (SQL)

- ISO standard way of using RDBMSs.
- Non-procedural (specifies what rather than how).
- Natural-language syntax resembling English sentences.
 - Keywords are case insensitive but CAPITALIZED as standard practice.
 - Data stored in database is case-sensitive.
 - Whitespace doesn't matter.
- Despite the name, has specifications for both a DDL and a DML.
- Frequently used for querying *even in non-relational databases*, so useful to know.

Dialects

- SQL comes in different dialects that support different features.
- Focusing in this class on MySQL, which can be used with the phpMyAdmin graphical user interface.



Data definition

Database creation

- Create a new database:
 - `CREATE DATABASE EstateAgent;`
 - (Can optionally specify character set that data will be stored in)
- Make it the current database in use:
 - `USE EstateAgent;`

Table creation

- Create a table to hold data:
 - `CREATE TABLE Staff (staffNo VARCHAR(4), fName VARCHAR(20), IName VARCHAR(20), position VARCHAR(20), sex VARCHAR(1), DOB VARCHAR (15), salary DECIMAL (5,2), branchNo VARCHAR(4));`
- Specify table name and list out attributes – name and type; optionally, properties about the attributes.

Data types

- BOOLEAN (three values: TRUE / FALSE / UNKNOWN).
 - Last one comes up due to null values.
- CHAR(#) – fixed length string (CHAR = CHAR(1)).
- VARCHAR(#) – strings not necessarily the same length in storage, w/ upper limit on length.
- BIT [VARYING](#) – binary strings.
- INT; DECIMAL – exact numeric values.
- FLOAT – approximate numeric values.
- DATE (YEAR/MONTH/DAY), TIME (HOUR/MINUTE/SECOND), TIMESTAMP (DATE + TIME), INTERVAL.
- CHARACTER LARGE OBJECT, BINARY LARGE OBJECT.

Data types: default values

- After specifying data type, can optionally set a default value:
 - `expirationDate DATE DEFAULT (CURRENT_DATE + INTERVAL 21 DAYS)`

Integrity constraints

- Allows you to specify properties about the data you would like the database to guarantee is true at all times.
- Five types:
 - Required data,
 - Domain constraints,
 - Entity integrity,
 - Referential integrity,
 - General constraints.

Integrity constraints: required data

- A field cannot contain NULL values (missing data).
- Specify this property during table creation by adding 'NOT NULL' to an attribute.

Integrity constraints: entity integrity

- Can specify a primary key, which uniquely identifies a record in a table (imposes uniqueness and non-null requirements).
 - Use 'PRIMARY KEY' keyword after an attribute.
 - Add 'PRIMARY KEY (col list...)' at the tail end of the attribute list.
 - Can also name this constraint by adding 'CONSTRAINT ConstraintName' before.
- Can specify non-primary key to be unique using UNIQUE keyword.

| Integrity constraints: referential integrity

- Foreign keys are special columns in one table that link its data to an entry in a second table by referencing that entry's primary key (from its own table).
 - Add 'FOREIGN KEY (colName) REFERENCES NameOfParentTable

Foreign key options

- What happens when the row my foreign key is referencing is deleted or its primary key is edited?
- Can optionally specify actions to take using 'ON {DELETE | UPDATE} {CASCADE | SET NULL | SET DEFAULT | NO ACTION}'
 - NO ACTION: Update/delete fails because dependents exist (default setting).
 - CASCADE UPDATE: This entry's foreign key is changed to match new value.
 - CASCADE DELETE: This entry is also deleted, possibly triggering further action if something is referencing this entry.
 - SET NULL: This entry's foreign key becomes null.
 - SET DEFAULT: This entry's foreign key is set to the column's default value (can be specified at table creation).

MySQL engines: InnoDB & MyISAM

- MySQL has different storage engines.
- Can set the storage engine for a table upon creation using “ENGINE = [INNODB, MYISAM]” flag.
- InnoDB is the default.
 - Implements foreign key feature and ACID properties for transactions.
 - Better than MyISAM for updating.
 - Worse than MyISAM for querying/searching.

Alter table definition

- Tables can be altered ('ALTER TABLE TableName') after the fact to:
 - ADD COLUMN colName type
 - DROP COLUMN colName
 - ADD CONSTRAINT constraintName ... (details)
 - DROP CONSTRAINT constraintName
 - ALTER colName {SET | DROP} DEFAULT

Delete table

- Delete a table using `DROP TABLE TableName`.



Data manipulation

Four main DML functions

- INSERT, SELECT, UPDATE, DELETE.
 - (Create, read, write, destroy.)

Populating the database with data

- Data can be inserted from a file stored locally to the database (such as from a text file called staff.txt):
- SQL load data statement:
 - E.g. `LOAD DATA LOCAL INFILE "staff.txt" INTO TABLE staff;`
- phpMyAdmin – use tools under the 'Import' tab.

Manual insertion of a row into a table

- `INSERT INTO Staff(staffNo, fName, LName, position, sex, DOB, salary, branchNo)
VALUES ('SG17', 'Joe', 'Brown', 'Manager', 'M', '12-09-90', 21000, 'B003')`

Update data in a table

- Give all staff a 3% pay increase.
 - UPDATE Staff
SET salary = salary*1.03
- Give all Managers a 5% pay cut.
 - UPDATE Staff
SET salary = salary * 0.95
WHERE position = 'Manager'

Delete data from a table

- `DELETE FROM Staff WHERE staffNo= 'SG17'`
 - Need WHERE to limit the scope of this action.
- `DELETE FROM Staff`
 - This deletes all rows in a table(!) while not deleting the table itself.

Components of a query to retrieve data

- Retrieval queries are composed from a limited set of clauses:
 - SELECT ... columns to appear in output
 - FROM ... table(s) to be used
 - WHERE ... filters rows by condition
 - GROUP BY ... forms groups of rows with same column value
 - HAVING ... filters groups subject to some condition
 - ORDER BY ... sorts the results

Example database

- Schema used in examples is for a fictional real estate company:
 - Staff (staffNo, fName, lName, position, sex, DOB, salary, branchNo);
 - Branch (branchNo, address, city);
 - Clients (clientNo, fName, lName, title, address, telNo, prefType);
 - Viewing (clientNo, propertyNo, viewdate, comment);
 - Property (propertyNo, address, type, rooms, rent, staffNo);

Simplest selection only uses SELECT and FROM

- *Task: List full details of all staff.*
 - SELECT staffNo, fName, lName, position, sex, DOB, salary, branchNo
FROM Staff
 - SELECT * FROM Staff

| staffNo | fName | lName | position | sex | DOB | salary | branchNo |
|---------|-------|-------|------------|-----|-----------|--------|----------|
| SA9 | Mary | Howe | Assistant | F | 19-Feb-70 | 10419 | B007 |
| SC 14 | David | Ford | Supervisor | M | 24-Mar-58 | 18900 | B003 |
| SG16 | Alan | Brown | Manager | M | 24-Jun-87 | 3675 | B002 |
| SG37 | Ann | Beech | Assistant | F | 10-Nov-60 | 13892 | B003 |
| SL21 | John | White | Manager | M | 01-Oct-45 | 3308 | B005 |
| SL41 | Julie | Lee | Assistant | F | 13-Jun-65 | 10419 | B005 |
| SW | Susan | Brand | Manager | F | 03-Jun-40 | 29172 | B003 |
| SG17 | Joe | Brown | Manager | M | 12-09-90 | 21000 | B003 |

Selection: filter out columns

- *Task: Produce a list of salaries for all staff, showing only staff number, first and last names, and salary.*
- `SELECT staffNo, fName, lName, salary
FROM Staff`

| staffNo | fName | lName | salary |
|---------|-------|-------|--------|
| SA9 | Mary | Howe | 10419 |
| SC 14 | David | Ford | 18900 |
| SG16 | Alan | Brown | 3675 |
| SG37 | Ann | Beech | 13892 |
| SL21 | John | White | 3308 |
| SL41 | Julie | Lee | 10419 |
| SW | Susan | Brand | 29172 |
| SG17 | Joe | Brown | 21000 |

Use DISTINCT keyword to remove duplicates from results

- *Task: List the property numbers of all properties that have been viewed.*
 - `SELECT DISTINCT propertyNo
FROM Viewing;`

SQL can apply mathematical operations to query results

- *Task: Produce a list of monthly salaries for all staff, showing staff number, first and last names, and salary details.*
 - `SELECT staffNo, fName, lName, salary / 12
FROM Staff;`
- Calculated columns are new columns and can be renamed:
 - `SELECT staffNo, fName, lName, salary / 12 AS monthlySal
FROM Staff;`

WHERE filtering: comparisons

- *Task: List all staff with a salary greater than 10,000.*
 - `SELECT staffNo, fName, lName, position, salary
FROM Staff
WHERE salary > 10000;`
- Other comparison operators: =, <>, <, >, <=, >=

WHERE filtering: compound conditions

- *Task: List addresses of all branch offices in London or Glasgow.*
 - `SELECT *`
`FROM Branch`
`WHERE city = 'London' OR city = 'Glasgow';`
- Other boolean operators: AND, OR, NOT

WHERE filtering: range

- *Task: List all staff with a salary between 20,000 and 30,000.*
 - `SELECT staffNo, fName, lName, position, salary
FROM Staff
WHERE salary BETWEEN 20000 AND 30000;`
- Ranges are inclusive: [20000, 30000]
- Can also use NOT BETWEEN
- Can implement using two `>=` or `<=` operators (multiple ways to skin a cat)

WHERE filtering: set membership

- *Task: List all managers and supervisors.*
 - `SELECT staffNo, fName, lName, position
FROM Staff
WHERE position IN ('Manager', 'Supervisor');`
- Can also use NOT IN
- Can implement using string of ORs

WHERE filtering: pattern-matching

- *Task: Find all properties with the string 'awr' in the street name.*
 - `SELECT propertyNo
FROM Property
WHERE street LIKE '%awr%'`
- Pattern is a regex with
 - %: wildcard; sequence of zero or more characters
 - _ (underscore): any single character.

WHERE filtering: NULL data

- *Task: List details of viewings where a comment has not been supplied.*
 - `SELECT clientNo, viewDate
FROM Viewing
WHERE comment IS NULL`
- Can also use `IS NOT NULL`

Sorting using ORDER BY

- *Task: List salaries for all staff, arranged in descending order of salary.*
 - `SELECT staffNo, fName, lName, salary
FROM Staff
ORDER BY salary DESC;`
- DESC = descending
ASC = ascending [default setting]

Sorting by multiple columns

- *Task: Produce an abbreviated list of properties in order of property type and rent.*
 - `SELECT propertyNo, type, rooms, rent
FROM Property
ORDER BY type, rent DESC;`

Aggregate functions

- Aggregation: calculation over a set of rows
- ISO standard defines five aggregate functions:
 - COUNT number of records (rows) in table incl. duplicates but not nulls
 - SUM total of values in specified column
 - AVG average of values in specified column
 - MIN smallest value in specified column
 - MAX largest values in specified column

Aggregate functions: details

- Operates on a set/group of values; returns single value
- COUNT, MIN, MAX for numeric and non-numeric fields
- SUM, AVG for numeric fields only
- Each aggregate function skips over nulls
- Include null values in count using COUNT(*)
- Remove duplicate values in count using DISTINCT keyword before column name

Using aggregate functions

- Aggregates can only be used in SELECT and HAVING clauses
- If an aggregate function is used in a SELECT clause and there is no GROUP BY clause, SELECT list cannot reference another column without using an aggregate function.
- Ex: illegal:
 - `SELECT staffNo, SUM(salary)`
`FROM Staff;`

Using COUNT(DISTINCT)

- *Task: How many different properties were viewed in May 2012?*
 - `SELECT COUNT(DISTINCT propertyNo) AS count
FROM Viewing
WHERE viewDate
BETWEEN '1-May-12'
AND '31-May-12';`

| |
|-------|
| count |
| 4 |

Using 'COUNT' and 'SUM'

- *Task: Find the number of managers and sum of their salaries.*
 - `SELECT COUNT(staffNo) AS count, SUM(salary) AS sum
FROM Staff
WHERE position ='Manager'`

| count | sum |
|-------|-------|
| 3 | 31927 |

Using 'MIN', 'MAX', 'AVG'

- *Task: Find minimum, maximum, and average staff salary.*
 - `SELECT MIN(salary) AS min, MAX(salary) AS max, AVG(salary) AS avg
FROM Staff`

| min | max | avg |
|------|-------|--------|
| 2921 | 25761 | 12452. |

GROUP BY: aggregation per category

- *Task: Find number of staff in each branch and their total salaries.*
 - `SELECT branchNo, COUNT(staffNo) AS count, SUM(salary) AS sum
FROM Staff
GROUP BY branchNo`

| branchNo | count | sum |
|----------|-------|-------|
| B002 | 1 | 3245 |
| B003 | 3 | 59537 |
| B005 | 2 | 13653 |
| B007 | 1 | 10732 |

GROUP BY details

- Calculates subtotals per category.
- SELECT and GROUP BY closely integrated:
 - One item in SELECT result per group in GROUP BY.
 - Column names in SELECT list *not* enclosed in an aggregate function must appear in the GROUP BY clause.
- SELECT clause may only contain:
 - Column names,
 - Aggregate functions,
 - Constants,
 - An expression involving combinations of the above.

HAVING vs. WHERE

- WHERE filters out rows *before* aggregate functions [over groups] are calculated.
- HAVING filters out group entries in the final aggregated table.

Using HAVING

- *Task: For each branch with more than 1 member of staff, find number of staff in each branch and sum of their salaries.*

- ```
SELECT branchNo, COUNT(staffNo) AS count, SUM(salary) AS sum
FROM Staff
GROUP BY branchNo
HAVING COUNT(staffNo) > 1
ORDER BY branchNo
LIMIT 0 , 30
```

| branchNo | count | sum   |
|----------|-------|-------|
| B003     | 3     | 59537 |
| B005     | 2     | 13653 |

# Subqueries

- SELECT queries can be nested within each other and within other DML functions (INSERT, UPDATE, DELETE).
- SELECT creates a temporary table: 0-D (scalar value), 1-D (single tuple/row or col), or 2-D (full table).
- Unlike aggregate functions, nested SELECTS can be inserted in WHERE / HAVING conditionals.

# Subquery using equality conditional

- *Task: List staff who work in branch at '163 Main St'.*
- ```
SELECT staffNo, fName, lName, position
FROM Staff
WHERE branchNo = (SELECT branchNo
                  FROM Branch
                  WHERE street = '163 Main St')
```


Produces 0-D (one-cell)
table with a single
branchNo

Final
result

| <u>staffNo</u> | fName | lName | position |
|----------------|-------|-------|------------|
| SC 14 | David | Ford | Supervisor |
| SG37 | Ann | Beech | Assistant |
| SW | Susan | Brand | Manager |

Subquery with aggregation

- *Task: List all staff whose salary is greater than the average salary.*
 - `SELECT staffNo, fName, lName, position, salary
FROM Staff
WHERE salary > (SELECT AVG(salary) FROM Staff)`



Can't use aggregate functions in a WHERE; use SELECT instead

Subquery using 'IN'

- *Task: List properties handled by staff at '163 Main St'.*
 - ```
SELECT propertyNo, street, city, postcode, type, rooms, rent
FROM Property
WHERE staffNo IN (
 SELECT staffNo
 FROM Staff
 WHERE branchNo = (
 SELECT branchNo
 FROM Branch
 WHERE street = '163 Main St'))
```
- Recall: “WHERE filtering: set membership” – 1-D results of SELECT queries constitute a set.

# 'ALL' and 'ANY' ('SOME')

- Used in comparisons to only select data for which the comparison is true for all or some values.
- Comparison will be performed for each value in a 1-D table (list of values).
- Edge case: empty result set -> ALL returns true, but ANY/SOME returns false.

# Subquery using 'ANY'/'SOME'

- *Task: Find staff whose salary is larger than salary of at least one member of staff at branch B003.*
- ```
SELECT staffNo, fName, lName, position, salary
FROM Staff
WHERE salary > SOME (
    SELECT salary
    FROM Staff
    WHERE branchNo = 'B003')
```

Subquery using 'ALL'

- *Task: Find staff whose salary is larger than salary of every member of staff at branch B003.*
- ```
SELECT staffNo, fName, lName, position, salary
FROM Staff
WHERE salary > ALL (
 SELECT salary
 FROM Staff
 WHERE branchNo = 'B003')
```

# 'EXISTS' and 'NOT EXISTS'

- Turns the result of a SELECT into true/false:
  - True if and only if there exists at least one row in result table returned by subquery.
  - False if subquery returns an empty result table.
- As (NOT) EXISTS checks only for existence or non-existence of rows in subquery result table, subquery can contain any number of columns.
  - Common for subqueries following (NOT) EXISTS to be of form: (SELECT \* ...)

# Using EXISTS on a subquery

- *Task: Find all staff who work in a London branch.*

- ```
SELECT staffNo, fName, lName, position
FROM Staff s
WHERE EXISTS (
    SELECT *
    FROM Branch b
    WHERE s.branchNo = b.branchNo AND city = 'London')
```

Temporary
table names

Table in larger
query is available
in subquery!

First condition matches staff to
branches, second filters for London

Multiple tables in FROM

- Specifying multiple comma-separated tables in a FROM clause will join them, giving you access to all columns and all combinations of rows from each table.
- Probably too many rows – need to filter using WHERE.

Two-table example

- *Task: List names of all clients who have viewed a property along with any comment supplied.*
 - `SELECT Client.clientNo, fName, lName, propertyNo, comment
FROM Client, Viewing
WHERE Client.clientNo = Viewing.clientNo`
- Can use aliases for table names:
 - `SELECT c.clientNo, fName, lName, propertyNo, comment
FROM Client c, Viewing v
WHERE c.clientNo = v.clientNo`

Other ways of joining

- Using FROM... JOIN... ON
 - SELECT ...
FROM Clients AS c
JOIN Viewings AS v ON c.clientNo = v.clientNo
 - Flexible; can use any condition after ON.
- FROM... JOIN... USING
 - SELECT ...
FROM Client JOIN Viewing USING clientNo
 - Succinct; assumes the tables both have columns with that name and you want them to be the same.

Three-table example

- *Task: For each branch, list staff who manage properties, including city in which branch is located and properties they manage.*
 - `SELECT b.branchNo, b.city, s.staffNo, fName, lName, propertyNo, type
FROM Branch b, Staff s, Property p
WHERE b.branchNo = s.branchNo AND s.staffNo = p.staffNo
ORDER BY b.branchNo, s.staffNo;`

A note on joins

- By default, joins are *inner* joins, meaning they only contain entries from the tables where the condition was fulfilled.
- Can also specify *outer* joins, which retain rows in one or both tables even if they didn't match with anything.
 - Missing attributes that would have been filled by the match(es) in the other table are filled with nulls.
 - Left join / left outer join preserves all rows from the first table.
 - Right join / right outer join preserves all rows from the second table.
 - Full outer join does both.

Outer join example

- Task: List branches and properties in same city, along with unmatched [branches | properties | branches or properties].
 - `SELECT b.*, p.*
FROM Branch b [LEFT | RIGHT | FULL] JOIN Property p
ON b.bCity = p.pCity;`

Set operations on tables: union, intersect, and difference ('EXCEPT')

- Tables must be *union compatible* (have attribute lists that have matching data types (domains)).
 - Union ('UNION') creates a table that contains rows in one or both tables.
 - Intersection ('INTERSECT') creates a table the only contains rows in both tables.
 - Difference ('EXCEPT') creates a table like the first table, minus any rows also found in the second table.
- Options:
 - Use keyword ALL to allow duplicates.
 - Use keyword CORRESPONDING BY (col list...) to restrict table / scope of operation to a (list of) column(s) – not supported in MySQL.

Union example

- *Task: List all cities where there is either a branch office or a property.*
 - (SELECT city FROM Branch)
UNION
(SELECT city FROM Property);
 - (SELECT * FROM Branch)
UNION CORRESPONDING BY city
(SELECT * FROM Property);

Intersect example

- *Task: List all cities where there is both a branch office and at least one property.*
 - (SELECT city FROM Branch)
INTERSECT
(SELECT city FROM Property);
 - (SELECT * FROM Branch)
INTERSECT CORRESPONDING BY city
(SELECT * FROM Property);

Except example

- *Task: List all cities where there is a branch office but no properties.*
 - (SELECT city FROM Branch)
EXCEPT
(SELECT city FROM Property);
 - (SELECT * FROM Branch)
EXCEPT CORRESPONDING BY city
(SELECT * FROM Property);



Advanced data definition

Integrity constraints: domain constraints

- Specify using CHECK after an attribute when creating table:
 - `CHECK (sex IN ('M', 'F'))`
 - `CHECK (branchNo IN (SELECT branchNo FROM Branch))`
- Create a new domain (custom data types that are defined separately from table construction and then usable):
 - `CREATE DOMAIN MyDomain AS CHAR DEFAULT 'M' CHECK (VALUE IN ('M', 'F'))`
 - `DROP DOMAIN [RESTRICT | CASCADE]`

Integrity constraints: general constraints

- Specify using CHECK again.
- Create a new assertion (custom data types that are defined separately from table construction and then usable):
 - `CREATE ASSERTION AssertionName CHECK (condition)`

Creating views

- Views allow users to get a facet of the database helpful to their specialized needs (see: subschemas).
 - You probably won't need them due to small database size.
- Create them from a query (SELECT):
 - `CREATE VIEW MyView [(opt. renamed cols list)] AS SELECT ...`
 - Can also specify how to handle scenario where an update to a row causes it to fall out of the view (if it has a WHERE clause)
 - `DROP VIEW MyView`
 - Can also specify how to handle deletion of views derived from this one

Views in SQL

- Can be convenient, simplifying, have security uses, etc.
- Limitations on the types of queries and updates that can be performed on a view because it is based on a SELECT.
- Performance overhead.
- Implemented either as a query that runs every time the view is manipulated, or as a separate temporary table that must be reconciled with the original(s) every so often.

Creating an index

- Creates a data structure allowing faster lookups of data.
- Makes ordered presentation or random access of a certain attribute instantaneous.
- `CREATE [UNIQUE] INDEX IndexName ON TableName (col or col list) [ASC | DESC]`

Transactions

- Can specify that a bundle of SQL statements should be committed or rolled back as a set (atomically).
- Use 'COMMIT'.
- Can specify beforehand how much the set of statements interact with each other.

Authorization

- SQL allows permissions to be set for users on parts of the database:
 - GRANT {ALL PRIVILEGES or a comma-separated list of SELECT, INSERT, UPDATE, DELETE, USAGE, REFERENCES} ON NameOfDatabaseTableEtc TO {PUBLIC or list of accounts}
 - REVOKE {privileges} ON ObjectName FROM {account(s) | PUBLIC}

Summary

- We've covered:
 - Data definition: creating tables with attributes.
 - Data manipulation: inserting, updating, and deleting rows; selection selection selection.
 - Data definition: constraints on data, views, security, etc.